

GreenFlag scoring and improvement report

Executive summary

GreenFlag is a **serious prize contender** because it solves a problem that is real, validated by FIA process documents, and unusually differentiated from the standard motorsport-AI field. The official challenge invites AI systems that improve how teams, drivers, or fans experience racing before, during, or after the race, with an emphasis on smarter decisions, explainability, and trust. GreenFlag fits that brief best when it is framed as a **pre-race, evidence-backed driver-enablement system** rather than an “AI that does FIA approval.” The public challenge page uses **technical execution, innovation, challenge fit, and real-world impact**; your brief uses **implementation & feasibility** as the fourth axis. On either framing, GreenFlag’s strongest advantages are originality, social impact, and explainability potential. ¹

The central strategic conclusion is simple: **GreenFlag should pitch itself as an audited drafting assistant for ASN/AWG review, not as an automated compliance oracle.** FIA Appendix L makes clear that the Certificate of Adaptations is issued after successful evaluation by the FIA Adaptations Working Group; the ASN must submit the request at least two months before competition; and only the adaptations specified in the certificate and appendices are authorized. The FIA’s Disability & Accessibility materials also state that the Vehicle Adaptation Guidelines are recommendations, not regulations, and that example dossiers are illustrative best practices rather than exhaustive or driver-specific advice. That regulatory structure is exactly why GreenFlag can win: it gives the project a clear human-in-the-loop role and a compelling auditability story. ²

My scoring view is that GreenFlag is currently around **18/20** if executed well, with a plausible path to first place if the demo visibly proves three things: every claim is source-backed, the system fails safely when evidence is weak, and an expert reviewer can inspect and override every red/yellow/green judgment. Without those elements, it is still highly competitive for **Most Innovative**, but more vulnerable on technical execution and feasibility. ³

Dimension	Score	Why it is strong	What keeps it from a perfect score
Technical execution	4.5/5	The stack is coherent: Docling/Granite Docling for parsing, Granite Vision for structured extraction, Granite 4.1 for drafting, Guardian for groundedness/BYOC checks, Langflow for visible orchestration, and Chroma/Postgres for retrieval + persistence.	The brief does not yet prove claim-level evidence binding, safe abstention, or reliable handling of uncontrolled phone photos.

Dimension	Score	Why it is strong	What keeps it from a perfect score
Innovation	5.0/5	This is one of the most original applications in the field: it attacks the accessibility/compliance bottleneck rather than yet another strategy, recap, or telemetry copilot.	Nothing major; the only risk is that the demo undersells the originality by looking like “document AI.”
Challenge fit	4.5/5	The challenge explicitly includes drivers, before/during/after race workflows, decision support, and explainability; GreenFlag is a strong “before the race” driver-access tool.	Some judges may initially see it as regulatory paperwork rather than a racing product unless the framing is very crisp.
Implementation and feasibility	4.0/5	The human-review posture is plausible, and official FIA materials strongly support an expert-reviewed workflow.	Medical/disability data, versioned regulations, and the fact that AWG retains decision authority all make “full automation” infeasible and risky.

The three highest-value changes are: **build a claim-by-claim audit layer**, **add an AWG reviewer mode with explicit override authority**, and **design the live demo around one safe success and one intentional failure case**. Those three moves improve technical execution, trust, and narrative simultaneously. ⁴

Regulatory ground truth

GreenFlag’s best pitch is grounded in a real FIA workflow. The FIA Disability & Accessibility Commission exists to facilitate safe participation of disabled people in motorsport, and the FIA states that a Certificate of Adaptations is required for adapted cars in international competitions or competitions that specifically require an FIA COA. The COA application must be sent at least two months before the competition. Appendix L further states that the certificate is issued after successful evaluation by the FIA Adaptations Working Group, is valid for the season if modifications remain unchanged, and must be presented at scrutineering. ⁵

That means GreenFlag should **never** imply that it “gets a driver approved.” The defensible claim is narrower and stronger: it **assembles a structured, source-backed draft dossier for expert review by the ASN and FIA AWG**. The FIA’s own vehicle adaptation guidance says the AWG has ultimate authority over evaluation of adapted vehicles, and the guidelines themselves “do not have regulatory value.” The FIA’s example technical dossiers are explicitly redacted/anonymized examples meant to inform drivers and adaptation experts, and the cover disclaimer says they do not provide advice about what is suitable for a particular vehicle or driver.

⁶

The regulatory corpus is also broader than Appendix L alone. The FIA Vehicle Adaptation Guidelines present nine major adaptation areas in their table of contents—throttle management, brake management, clutch system, steering assembly, gearshift and gearbox, seat and driver restraints, headrest and cockpit

environment, driver equipment, and chassis modification—which makes GreenFlag’s nine-domain scorecard a credible organizing structure if it is clearly labeled as a **best-practice review frame**, not an FIA regulation in itself. The same guidance stresses that adaptations must be safe, individualized, and must not provide a sporting advantage. 7

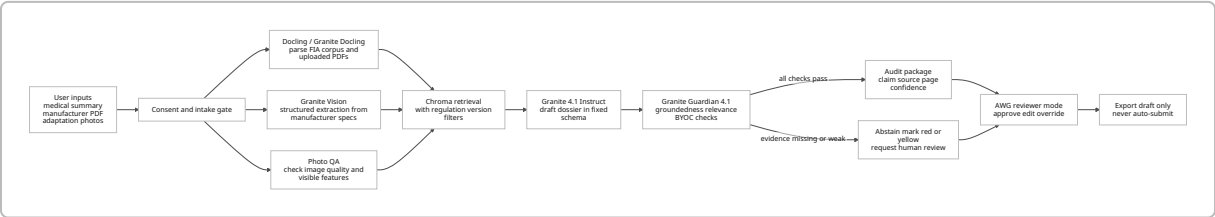
The standards cross-references in your brief also make technical sense. FIA materials publicly reference 8855-2021 and 8862-2009 seat standards, 8858-2010 frontal head restraint standards, and 8873-2018 karting high seats; the adaptation guidance specifically recommends certain restraint and seat combinations and notes that some adapted protective clothing modifications must be done by the manufacturer and accompanied by appropriate certification. In other words, GreenFlag’s value is not in “knowing one rule,” but in combining Appendix L, adaptation guidance, technical lists, and discipline-specific examples into a single reviewable draft. 8

This is also why the market narrative is stronger than it first appears. FIA’s disability page emphasizes inclusion in motorsport, while the WHO estimates that about **1.3 billion people**, roughly **1 in 6**, experience significant disability globally. Team BRIT describes itself as the world’s only competitive team of all-disabled racing drivers, and Mission Motorsport says its programs have reached more than 2,000 veterans and service leavers. GreenFlag does not need an inflated TAM claim to feel real; the community and the process bottleneck are already well evidenced. 9

Technical execution assessment

The stack you described is stronger than a generic “RAG over racing rules” build because each component has a plausible, distinct role. Docling and Granite Docling are built for structured document conversion; Granite Vision 4.1 is explicitly positioned for structured extraction tasks such as tables and charts; Granite 4.1 8B is a long-context instruct model suitable for extraction, QA, and RAG workflows; Guardian 4.1 can judge groundedness, relevance, and arbitrary BYOC criteria; Langflow gives a visible, serializable execution graph; Chroma supports metadata-aware retrieval; and PostgreSQL gives you the security and audit primitives you actually need for sensitive dossier workflows. 10

The most important architecture change is to make the system **evidence-first rather than generation-first**. Granite 4.1 has a 131,072-token sequence length, but that does not mean the best strategy is to dump the whole corpus into the prompt. For regulatory drafting, the better pattern is: parse the standards and guidance into structured objects, retrieve only the most relevant chunks with metadata filters, draft into a constrained template, and refuse any sentence that lacks a source anchor. Long context expands possibility; it does not remove the need for retrieval discipline. 11



The strongest verified use of each stack element is summarized below.

Stack element	Best verified use in GreenFlag	Main limitation GreenFlag must design around
Docling / Granite Docling	Parse FIA PDFs and manufacturer documents into structured representations; export tables, figures, and images for citation bundles.	IBM warns outputs can be inaccurate or hallucinatory; it should be treated as a converter, not a policy authority.
Granite Vision 4.1 4B	Extract structured values from charts, tables, and document-like images, especially manufacturer spec PDFs.	The model card says it is designed for structured extraction tasks and may not generalize well to open-ended VLM tasks; phone photos are the weak link.
Granite 4.1 8B Instruct	Produce a tightly templated draft dossier after retrieval and evidence assembly.	IBM notes it can still produce inaccurate, biased, or unsafe outputs; unconstrained free-form drafting is risky.
Granite Guardian 4.1 8B	Gate outputs using groundedness, answer relevance, and BYOC rules such as “every compliance claim must cite a source.”	Guardian is a yes/no judge, not proof of correctness; IBM says it must be used in prescribed scoring mode and custom criteria require testing.
Langflow	Make the orchestration visible to judges and reviewers; serialize the flow for reproducibility.	A visual graph helps only if each node corresponds to a real artifact on-screen.
Chroma	Store embeddings plus metadata and filter retrieval by standard, version, discipline, and page.	Chroma’s in-memory mode is non-persistent; production requires explicit persistent or client-server setup.
PostgreSQL	Store driver records, audit objects, source hashes, and reviewer decisions in <code>jsonb</code> ; enforce access control and logging.	Sensitive data handling is on you: encryption, RLS, and audit tables must be intentionally configured.

This architecture table is grounded in the official Docling docs, IBM model cards, Langflow docs, Chroma docs, and PostgreSQL docs. Docling supports exporting tables, figure/page images, and VLM pipelines; Granite Vision documents structured chart/table extraction and warns against assuming open-ended task generalization; Granite 4.1 states intended use for RAG and agentic workflows; Guardian documents groundedness/answer-relevance/BYOC scoring; Chroma supports metadata-aware retrieval and persistent deployments; and PostgreSQL documents `jsonb`, row-level security, encryption, and audit-trigger patterns. ¹²

The single largest technical vulnerability is **image extraction from free-form adaptation photos**. Granite Vision 4.1 is impressive, but IBM’s own card says it is designed for structured extraction tasks, uses English instructions only, and may not generalize well to open-ended VLM tasks. That is fine for manufacturer spec PDFs; it is not enough to justify blind trust in a handheld phone photo of a custom steering or restraint setup. GreenFlag should therefore treat photos as **supporting evidence for reviewer inspection**, plus a limited set of controlled checks, unless the capture process is standardized. ¹³

Auditability, image reliability, and governance

GreenFlag's winning move is not "better generation." It is **better auditability**. The public challenge explicitly talks about explainability and trust, and the FIA process is inherently review-heavy. Guardian 4.1 is especially useful here because IBM documents both pre-baked hallucination checks for RAG pipelines and BYOC checks for arbitrary user-defined rules. That means GreenFlag can implement concrete rules like: *every red/yellow/green domain label must cite at least one source chunk; any unsupported sentence must be removed or rewritten as unknown; if a claim relies on a photo, that photo must be visible in the reviewer panel.* That is a much stronger story than merely saying "we used Guardian for safety." ¹⁴

At the same time, Guardian has real limits. IBM says Guardian must be used in its prescribed scoring mode; reasoning traces may be unsafe or unfaithful; the model is trained and tested only on English; and BYOC criteria still require testing. The right operational role is therefore **gating and observability**, not final judgment. In the demo, show the verdicts and evidence anchors, not the model's hidden reasoning. ¹⁵

The governance posture also has to be tighter than a normal hackathon app because GreenFlag touches **health and disability data**. Official UK GDPR/ICO guidance treats health data as special-category data, and disability may be special-category data insofar as it reveals health information. ICO guidance also states that special-category data cannot be used for solely automated decision-making with legal or similarly significant effects unless strict conditions are met, and DPIAs are required where processing is likely to result in high risk. A system that drafts racing-adaptation dossiers from disability and medical context should therefore be designed as a **human-reviewed assistant**, with explicit consent, minimization, encryption, retention limits, and documented impact assessment. ¹⁶

PostgreSQL can support the right controls if used deliberately. The official docs note support for encryption options, `jsonb` storage, row-level security policies, and audit-trigger patterns. In practice, that means GreenFlag should store structured evidence objects in `jsonb`, enable RLS by tenant or reviewer role, encrypt sensitive fields, and create immutable audit tables for draft generation, edit decisions, and export events. Without that, the "Driver Passport" idea becomes a governance liability. ¹⁷

A risk-led design for GreenFlag looks like this:

Risk area	Why it matters	What GreenFlag should implement
Unsupported compliance language	A polished dossier that invents regulatory support is worse than no dossier.	Claim-level source anchors; no unsupported sentence may survive export.
Free-form photo misread	Uncontrolled cockpit/adaptation photos are the least validated part of the stack.	Standardized capture instructions, quality checks, fiducial/ruler card, manual confirmation, fail-closed behavior.
Regulation drift	FIA guidance is versioned and one document is explicitly a living document.	Versioned corpus, visible regulation date on every export, invalidate old caches after updates.

Risk area	Why it matters	What GreenFlag should implement
Automated overreach	AWG, not GreenFlag, decides.	"Draft only" watermark, reviewer sign-off gates, no auto-submit.
Medical/disability privacy	Special-category data raises legal and trust risk.	Consent, minimization, encryption, RLS, audit logs, deletion schedule, DPIA.
False certainty in safeguards	Guardian is useful but not infallible.	Guardian as gate + spot-checker, not sole arbiter; reviewer override required.

The audit table below is the kind of artifact that would materially strengthen the demo.

Domain	Draft statement	Evidence anchor	Model confidence	Guardian verdict	Reviewer action
Submission timing	COA request must go via the ASN at least two months before the competition.	Appendix L, Article 18.3, source chunk ID ap126_p29_art18_3_a	0.98	Grounded = yes	Accept
Authorization scope	Only the modifications specified in the certificate and appendices are authorized.	Appendix L, Article 18.3, source chunk ID ap126_p28_art18_3_b	0.97	Grounded = yes	Accept
Guideline status	Vehicle Adaptation Guidelines are best-practice guidance, not binding regulation.	Vehicle Adaptation Guidelines intro, source chunk ID vag_p6_intro	0.96	Grounded = yes	Add advisory footnote

Domain	Draft statement	Evidence anchor	Model confidence	Guardian verdict	Reviewer action
Template reuse limitation	Example technical dossiers are examples, not advice for a particular vehicle or driver.	Example Technical Dossier disclaimer, source chunk ID <code>exa_p0_disclaimer</code>	0.95	Grounded = yes	Keep disclaimer

Those sample rows are based directly on Appendix L, the Vehicle Adaptation Guidelines, and the FIA example technical dossier disclaimer. ¹⁸

The most valuable product feature to add is therefore an **AWG reviewer mode** with five capabilities: side-by-side source viewer, per-domain approve/edit/override controls, automatic diffing between machine draft and reviewer draft, immutable decision logs, and an export footer that records corpus version and reviewer identity. That feature aligns with FIA authority structure, reduces privacy/regulatory anxiety, and gives judges something concrete to trust. ¹⁹

Demo, narrative, and competitive position

GreenFlag’s demo should be designed around the challenge’s published emphasis on smarter decisions, explainability, and trust. The public challenge page explicitly asks builders to show how AI can improve racing experiences before, during, or after the race, and to use explainability to build trust. GreenFlag already fits the “before the race” part; the missing piece is to make **explainability visible** enough that judges can feel it in under a minute. ²⁰

The narrative framing should also be sharpened. Start with the barrier, not the paperwork: **“Most racing AI helps drivers go faster once they’re already on track. GreenFlag helps disabled drivers get cleared to race.”** That line is cleaner than a longer detour into regulation. If you want an IBM-brand bridge, use it briefly: IBM publicly highlights a long history of disability inclusion, has a recent Inclusive Brains research partnership, and already uses watsonx with Scuderia Ferrari to turn racing data into explainable fan features. That gives GreenFlag strong adjacency without making the pitch feel like corporate theater. ²¹

The live demo should include the following elements, in this order:

Demo element	Why judges will care	What to show
One fast “happy path” draft	Proves the core magic in under 90 seconds.	Short medical summary + manufacturer PDF + pre-approved clean photo set; dossier appears with domain scorecard.
Live audit screen	Converts “cool output” into “trustworthy system.”	Click a claim and show source chunk, page number, page image, and Guardian result.

Demo element	Why judges will care	What to show
Intentional failure case	Demonstrates safety and honesty.	Use a blurry or incomplete photo/spec so one or more domains turn yellow/red with “insufficient evidence.”
AWG reviewer mode	Shows human-in-the-loop feasibility.	Reviewer accepts one claim, edits one, and overrides one.
Timing panel	Makes the 90-second workflow believable.	Stage timings for parse, retrieve, draft, guard, and export.
Draft-only watermark	Prevents overclaiming.	“DRAFT — REQUIRES ASN/AWG REVIEW” on every page and on export.

The phone-photo step deserves special handling in the demo. Because Granite Vision’s verified strengths are more document-centric than free-form-photo-centric, the best version of the live demo is to use **controlled photo capture**: front/side angle, good lighting, checklist overlay, and an obvious quality indicator. If the image fails quality checks, the system should explicitly downgrade confidence and ask for another photo. That safe failure is not a weakness; it is one of the most convincing things GreenFlag can show. ²²

On competitors, GreenFlag’s relative position is strong but not guaranteed. The comparison below is **scenario-based**, using the competitor descriptions in your supplied field brief rather than independently verified public submissions. [filecite:turn0file0](#) [filecite:turn0file1](#)

Head-to-head	Current edge	Why
GreenFlag vs PitWall	Slight GreenFlag edge, but conditional	PitWall likely wins on instant comprehensibility and cinematic energy; GreenFlag wins on originality, social impact, and a stronger “AI beyond speed” story. If GreenFlag’s audit layer is invisible, PitWall can win the room.
GreenFlag vs RaceMind AI	GreenFlag	RaceMind sounds competent but closer to common telemetry/commentary territory; GreenFlag is more differentiated and more emotionally memorable.
GreenFlag vs AI Race Strategist	GreenFlag	Strategy assistants are likely crowded in this field; GreenFlag avoids the clone zone.
GreenFlag vs strong AI Race Steward	Toss-up	A steward project could score very highly on explainability if it shows video, rule citations, and reversible reasoning; GreenFlag still has the stronger accessibility mission and less crowded problem space.

My prize-likelihood view is therefore:

Prize	Probability	Rationale
First place	55%	Very plausible if the demo visibly proves auditability and safe failure; less likely if it looks like a fancy PDF generator.
Runner up	75%	High floor because the problem is real, differentiated, and aligns with explainability and impact.
Best Use of Technology	58%	Strong IBM-stack story, but only if each IBM component has an obvious on-screen job rather than a slide-deck mention.
Most Innovative	82%	This is the most natural category for GreenFlag because the problem framing is genuinely unusual in the field.
Grand prize	28%	Possible, but dependent on the June field and on whether GreenFlag's execution looks product-grade rather than hackathon-grade.

Action roadmap and verdict

The fastest path to a winning build is not a major rewrite. It is a **tightening of evidence, review, and demo choreography**.

Priority	Change	Why it matters most	Effort
Highest	Build claim-level evidence binding	Directly raises technical execution and trust; every sentence in the draft should carry source IDs, page refs, and confidence.	Medium
Highest	Ship AWG reviewer mode	Converts regulatory risk into a strength by making expert oversight visible and indispensable.	Medium
Highest	Add one intentional failure demo	Proves the system is honest under uncertainty; judges remember this.	Low
High	Standardize photo intake	Reduces the biggest model-risk area and makes the output more believable.	Medium
High	Add regulation/version banner	Prevents “stale rules” objections and conveys professionalism.	Low
Medium	Split templates by discipline	FIA examples span endurance, GT3, rally, and rally-raid; one generic template is less credible than discipline-aware variants.	Medium

A compressed implementation plan can fit into a short sprint:

Sprint window	Milestone	Delivery artifact
Days one through three	Evidence schema and source registry	<code>source_id</code> , document hash, version, page, block/image ref, confidence stored in Postgres <code>jsonb</code>
Days four through six	Guardian gates and abstention logic	BYOC checks for citation presence, groundedness, unsupported-claim rejection, and fail-closed export rules
Days seven through nine	AWG reviewer mode	Review queue, approve/edit/override, diff viewer, immutable audit log
Days ten through eleven	Photo-capture and failure-case flow	Quality scoring, retry prompt, red/yellow logic for weak evidence
Days twelve through fourteen	Demo polish and timing instrumentation	Stage timers, “happy path” script, “failure path” script, final video

The honest verdict is favorable. GreenFlag is not strongest when described as “AI writes the FIA dossier for you.” It is strongest when described as **an evidence-backed drafting and review system that shortens preparation time while preserving expert authority, safety, and auditability**. Under that framing, it is one of the best-aligned “AI Beyond the Finish Line” concepts I have seen: it addresses drivers before the race, directly serves an underserved community, and can make explainability the centerpiece rather than an afterthought. If the live demo includes the audit screen, the intentional failure case, and AWG reviewer mode, GreenFlag is a credible **first-place contender** and the most natural candidate for **Most Innovative**. Without those changes, it still looks strong, but more likely as **Runner Up + Most Innovative** than as the outright winner. ²³

Open questions and limitations

This report is grounded in official FIA, IBM, Docling, Langflow, Chroma, PostgreSQL, and privacy-regulator sources, plus the GreenFlag and competitor briefs you provided. I did **not** independently review a GreenFlag code repository, live demo, or internal architecture diagrams, so technical conclusions are about **design plausibility and competition strategy**, not implementation verification. I also did not rely on every field/TAM claim in the brief where a strong public primary source was not readily available; where public evidence was incomplete, I treated those points as secondary rather than load-bearing.

¹ ³ ¹⁴ ²⁰ ²³ <https://ibmskillsbuildchallenge.bemyapp.com/>

<https://ibmskillsbuildchallenge.bemyapp.com/>

² ¹⁸ https://www.fia.com/system/files/documents/appendix_l_2026_public_le_26_mars_2026.pdf

https://www.fia.com/system/files/documents/appendix_l_2026_public_le_26_mars_2026.pdf

⁴ ¹⁵ <https://huggingface.co/ibm-granite/granite-guardian-4.1-8b>

<https://huggingface.co/ibm-granite/granite-guardian-4.1-8b>

⁵ ⁹ <https://www.fia.com/disability-accessibility>

<https://www.fia.com/disability-accessibility>

6 7 19 https://api.fia.com/sites/default/files/vehicle_adaptation_guidelines.pdf
https://api.fia.com/sites/default/files/vehicle_adaptation_guidelines.pdf

8 https://api.fia.com/sites/default/files/fia_standard_guidelines_-_seats.pdf
https://api.fia.com/sites/default/files/fia_standard_guidelines_-_seats.pdf

10 <https://huggingface.co/ibm-granite/granite-docling-258M>
<https://huggingface.co/ibm-granite/granite-docling-258M>

11 <https://huggingface.co/ibm-granite/granite-4.1-8b>
<https://huggingface.co/ibm-granite/granite-4.1-8b>

12 https://docling-project.github.io/docling/examples/export_tables/
https://docling-project.github.io/docling/examples/export_tables/

13 22 <https://huggingface.co/ibm-granite/granite-vision-4.1-4b>
<https://huggingface.co/ibm-granite/granite-vision-4.1-4b>

16 <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/a-guide-to-lawful-basis/special-category-data/>
<https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/a-guide-to-lawful-basis/special-category-data/>

17 <https://www.postgresql.org/docs/current/encryption-options.html>
<https://www.postgresql.org/docs/current/encryption-options.html>

21 <https://www.ibm.com/history/equal-workforce>
<https://www.ibm.com/history/equal-workforce>