

Adversarial Technical Review of PIT WALL

Executive summary

Your memo describes **PIT WALL** as a motorsport decision copilot that ingests telemetry CSVs, FIA timing PDFs, and written debrief text, then forecasts next-session lap pace and produces a short natural-language explanation on an IBM Granite stack. [filecite](#)[turn0file1](#)

My bottom-line judgment is this: **the concept is technically plausible only in a narrowed form, but the current public Granite TTM evidence does not justify a strong claim that raw high-frequency motorsport telemetry is already a good fit for the shipping IBM Granite TimeSeries stack.** IBM's public Granite-TimeSeries-TTM-R2 card says the open release supports minutely and hourly data; the current r2.1 release adds daily and weekly support, but still frames TTM as a family of small, focused forecasting models selected by context length and horizon rather than as a universal streaming telemetry model. The public R2 model is only about **805k parameters**, and IBM positions it as lightweight, laptop-friendly, and suitable for exogenous/control-variable integration. ¹

The **benchmark story is real but weaker than the pitch implies.** The NeurIPS paper does show strong average gains over several well-known baselines, including Moirai, TimesFM, Chronos, Lag-Llama, GPT4TS, and Moment in selected settings. But those comparisons are not fully apples-to-apples: some baselines are evaluated on **sliding-window** zero-shot tests while others are evaluated on **last-window only**; some runtime numbers are taken under a special setup with **batch size 32, context length 512, horizon 96, one channel only**, and probabilistic models are timed with **num_samples = 1** to get their fastest path; and the benchmark tables I reviewed report **point estimates** rather than confidence intervals or statistical significance tests. ²

For an investor or a technically sophisticated customer, the most honest conclusion is: **TTM itself is unlikely to be the latency bottleneck, but the public Granite TTM release is not yet a clean off-the-shelf fit for raw motorsport telemetry forecasting at the granularity implied by the memo.** The likely path to a credible one-week prototype is not "feed raw race telemetry to Granite TTM and win," but rather "engineer motorsport-aware aggregates, use TTM or a classical baseline on those aggregates, and let the LLM explain outputs." That is still viable, but it is a narrower and less novel claim than the memo suggests. ³

What the public IBM Granite TTM stack really supports

IBM's current public Granite-TimeSeries-TTM-R2 model card describes TTM as a family of **focused pretrained multivariate forecasters** rather than one massive general model. The public card says **TTM r2** was trained on roughly **700M samples**, **r2.1** increases that to roughly **1B**, and the released R2 model card exposes automatic model selection by **context length**, **prediction length**, and frequency preferences. The same card says the family supports **multivariate forecasting**, **exogenous/control-variable integration**,

and **channel-mixing** during fine-tuning, which is directly relevant to a race-engineering use case that wants to combine telemetry with timing or contextual variables. ¹

The first serious technical problem for the PIT WALL pitch is **frequency support**. The public Granite-TimeSeries-TTM-R2 card says the open release supports **minutely and hourly** resolutions, while **r2.1** adds **daily and weekly** support. The underlying `get_model()` code reinforces that framing: when the requested frequency maps to “daily or longer,” the loader explicitly filters to **r2.1** models, and unsupported frequencies can lead to a failure path that requires either rolling, zero-padding, or random initialization instead of a clean pretrained match. The same model card warns that **upsampling** or **prepending zeros** to force a context fit is **not recommended**, and that rolling forecasts **beyond 2× prediction length** are not recommended either. ⁴

That matters because the **research paper and the shipping public release are not the same thing**. The NeurIPS paper says TTM was pretrained on data spanning **seconds to days**, and the appendix lists pretraining sets that include **4-second** solar and wind series plus several downsampled multi-resolution datasets. The evaluation appendix also includes D2 datasets at **10 seconds** and **2 minutes**. But the public Granite model card explicitly says the released open model differs slightly from the research-paper pretraining data because IBM restricted the release to datasets with clear commercial-use licensing; IBM also notes that this can lead to **minor performance differences** relative to the paper. In other words, the paper gives evidence that the TTM research line can touch faster signals, while the shipping enterprise-friendly card narrows the usable envelope. ⁵

There is a second structural mismatch: the paper’s own limitations section says TTM is **sensitive to context length**, requires different variants for different context settings, and currently supports **point forecasting only**, unlike some competing models that support probabilistic forecasting. For a race-engineering workflow, that means the public TTM story is stronger for “give me a point estimate of the next value or short horizon” than for “give me calibrated uncertainty for strategy tradeoffs.” ⁶

Why the benchmark story is weaker than the headline

The paper’s headline claims are not fabricated. In the main paper, Table 1 reports average zero-shot forecast improvements of TTM over **Moirai** and **TimesFM**; Table 2 reports zero-shot gains over **Chronos** and **Lag-Llama** on a different protocol; Table 4 reports 5% few-shot gains over **Moment** and **GPT4TS**; and Table 5 reports modest but positive gains over **Moment**, **GPT4TS**, and **Time-LLM** in full-shot head probing. IBM’s public Granite model card repeats the broad claim that TTM outperforms several popular larger baselines while remaining tiny and CPU-friendly. ⁷

The problem is that the benchmark package is **heterogeneous**. In the appendix, Table 10 documents that different baselines were sourced in different ways: some results were generated from **official implementations**, some were taken from **other papers**, and some were used only for specific tables. The paper also changes the evaluation protocol across baseline groups: the Moirai/TimesFM comparison uses **sliding-window** results across horizons **96, 192, 336, 720**, while the Chronos/Lag-Llama comparison uses **last-window** results and shorter horizons **24, 48, 60, 96, 192** because those models recommend or report shorter horizons. That makes the overall story informative, but it is not a single unified benchmark with one protocol and one inference regime for all competitors. ⁸

The runtime comparison is especially favorable to TTM. Appendix D.3 says Table 3 measures runtime with **context length 512, forecast length 96, batch size 32, and one channel only** to keep the comparison fair for models that are fundamentally univariate. It also says probabilistic models such as Chronos and Lag-Llama were timed with **num_samples = 1** to compute the fastest possible runtime, even though the paper uses larger sample counts for their forecast-accuracy comparisons. Those conditions are useful for a controlled benchmark, but they are materially easier than a real motorsport multivariate pipeline that includes multiple telemetry channels, feature extraction, context assembly, and narrative generation. ⁹

The tables also show that TTM is **not dominant on every dataset/horizon pair**. In the full sliding-window zero-shot appendix table, Moirai beats or ties TTM on some cells, including examples such as **ETTH1 at horizon 720** and some **Electricity** settings. In the last-window table against Chronos and Lag-Llama, there are settings where Chronos or Lag-Llama are better, such as **Weather at horizon 24, Electricity at horizon 24**, and parts of the **Traffic** table. So the strongest defensible statement is “TTM wins strongly on average in the authors’ chosen summaries,” not “TTM is universally best.” ¹⁰

Finally, the benchmark presentation is **not statistically rigorous enough for sweeping claims**. In the main tables and setup sections I reviewed, the paper reports averaged **MSE**, parameter counts, and runtime, but I did not see confidence intervals, repeated-seed variance, or formal statistical significance testing attached to the headline comparisons. For a startup pitch, that means you can honestly cite the tables, but you should not overstate them as statistically settled proof that TTM will generalize into motorsport telemetry. ¹¹

Fit for motorsport telemetry and one-week prototype feasibility

For the use case in your memo, the key question is not “Is TTM fast?” but “**Can the public Granite TTM release ingest the representation you actually have?**” On that point, I am skeptical. The public IBM Granite card is clear that released TTM models are organized around supported frequencies and minimum context lengths, and the shipping open release is framed around minutely, hourly, daily, and weekly resolutions rather than arbitrary high-frequency sensor streams. If your intended pipeline relies on raw or near-raw motorsport telemetry at very fine granularity, then the public enterprise-friendly TTM release is **outside its stated sweet spot**. ¹²

That said, your memo can still be made technically credible if you **narrow the forecasting object**. A one-week prototype can be believable if the target becomes something like: **predict next-lap or next-session pace from engineered aggregates** such as sector deltas, throttle/brake distributions, corner-entry/exit summaries, tire-age proxies, and exogenous race-context variables extracted from timing sheets and debrief text. That reframing aligns much better with what TTM is actually good at: compact multivariate forecasting with exogenous/control inputs once the raw telemetry has already been transformed into a structured time-series representation. The paper’s D2 experiments and model card both support the exogenous-variable story. ¹³

There is even a subtle positive signal here. The NeurIPS appendix shows that the research version of TTM handled **10-second** evaluation datasets with exogenous/control variables and that the broader research pretraining corpus included **4-second** data, so the underlying TTM line is not philosophically incompatible with faster telemetry-like signals. The obstacle is that **the public Granite release you can honestly pitch today is narrower than the research evidence**. That is exactly the kind of distinction a technical reviewer or diligence team will catch. ¹⁴

So my adversarial verdict on feasibility is: **a constrained prototype is medium-feasibility; a raw-telemetry foundation-model claim is low-feasibility**. If you pitch the latter, you risk being falsified quickly by anyone who reads the IBM model card. If you pitch the former, you have a defensible path. ¹⁵

Latency and deployment realism

On pure forecasting latency, TTM looks genuinely strong. In Table 3 of the NeurIPS paper, **TTM-B** is reported at roughly **4.7 ms GPU inference time per batch** and **0.01 s CPU time per batch** under the paper's runtime setup. A later NOMS 2026 paper evaluating **TTM on ISP telemetry** reports CPU-only inference latency of **under 0.05 s per 100 points on a MacBook Pro**, which is useful evidence that TTM-like telemetry forecasting can be operationally cheap on commodity hardware. If PIT WALL's forecasting block is well designed, the forecaster itself is unlikely to be the bottleneck. ¹⁶

The harder SLA question is the **LLM narration layer**. IBM's Granite 4.1 documentation and Hugging Face model card say Granite 4.1 is available in **3B, 8B, and 30B**, with the **8B** variant positioned as the balanced general-purpose model. IBM's official materials show how to run it via **Transformers, vLLM, and SGLang**, and IBM Research says Granite 4.1 is optimized for local runtimes including **llama.cpp**. The official IBM materials I reviewed emphasize "predictable latency" and local deployability, but they do **not** publish local token-per-second benchmarks for an **M2 Mac** or **RTX 4060 laptop**. ¹⁷

There are still two practical signals. First, Ollama's Granite 4.1 library lists the **8B** package at about **5.3GB** with a **128K context window**, which means a quantized local deployment is at least memory-feasible on midrange consumer hardware. Second, Artificial Analysis reports a hosted API speed of about **118 output tokens/sec** with **0.79 s** time-to-first-token for Granite 4.1 8B through CoreWeave, which is an upper-bound style reference point for what well-provisioned serving can do. But because IBM does not publish local M2/4060 measurements in the official materials I reviewed, any claim that "the full local pipeline will finish in under 60 seconds" is still an **engineering estimate**, not a sourced benchmark. ¹⁸

My practical inference is therefore narrow: **a sub-60-second demo is plausible if the TTM block is small, PDF extraction is lightweight, and the explanation is short and tightly templated**; it is **not yet proven** for a general free-form "telemetry + PDF + debrief -> rich narrative" local pipeline using Granite 4.1 8B on commodity laptops, because IBM does not publish those local speed numbers. If the SLA matters commercially, a smaller narration model such as **Granite 4.1 3B** or a partially templated output layer is the safer design choice. ¹⁹

Prior art and defensibility

I did **not** find a motorsport-specific IBM Granite TTM case study in the official IBM Research, Granite community cookbook, or Hugging Face Spaces searches I ran. That is important because it means the memo cannot honestly say "IBM Granite already proves this on motorsport telemetry." At best, it can say the IBM stack is promising and adjacent-domain evidence exists.

But the related-work landscape is already crowded enough to weaken the novelty claim. In **sim racing**, Hojaji and coauthors used telemetry from **174 participants** across **1,327 laps**, logged with **MoTeC i2 Pro**, and showed that machine learning could predict lap-performance groupings with very high accuracy while surfacing key driving variables such as average speed, lateral acceleration, steering angle, and lane

deviation. A related Springer chapter from 2022 argues that telemetry-rich sim-racing performance prediction was already an open ML problem and used public **MoTeC** datasets for it. Separately, a 2023 thesis built an **AI-based race strategy assistant** that predicts pit-stop timing and tire choices from race-state and telemetry-derived inputs. ²⁰

In adjacent foundation-model work, the idea is also no longer novel. A 2025 arXiv paper applied **Chronos** to **car-following behavior analysis**, showing that a time-series foundation model can outperform traditional traffic models and approach or beat deep-learning baselines in an automotive behavior setting. A 2026 NOMS paper evaluated **TTM on ISP telemetry**, explicitly positioning TTM as viable for real telemetry forecasting with CPU-only deployment. A 2025 Computer Networks paper explored TSFMs, including TTM, for **network monitoring and video QoE inference** from coarse telemetry-like signals. So the general idea of “foundation models for operational telemetry” is already established outside motorsport. ²¹

That means PIT WALL’s defensibility is **not** the model choice. It would have to come from a **data moat** and an **evaluation moat**: proprietary race-weekend telemetry, trustworthy labels for “next-session pace,” careful out-of-event validation, useful extraction of contextual documents, and first-class workflow integration for engineers and analysts. If you do not have those, then PIT WALL is mostly a thin application layer over publicly available models and already-published ideas. ²²

Investment-grade verdict

If I were writing the diligence note, I would score the concept this way.

Technical feasibility: medium, but only after narrowing the problem. A prototype that forecasts from **engineered lap/session features plus exogenous race context** is credible. A prototype that claims to directly operationalize the public Granite TTM release over raw fine-grained motorsport telemetry is not. ²³

Novelty: low to medium. There is enough prior art in sim-racing telemetry, race-strategy ML, automotive foundation-model forecasting, and telemetry-domain TSFM deployment that “AI for motorsport telemetry forecasting” is not itself novel. The only potentially novel layer is a proprietary, production-ready motorsport workflow with strong evaluation and valuable data access. ²⁴

Benchmark credibility: medium at best. The TTM paper is strong evidence that IBM has a serious lightweight forecaster, but it does **not** justify a blanket claim that TTM will dominate on PIT WALL’s specific data. The evaluation protocols vary, the runtime setup is simplified, and the shipping commercial-friendly release differs from the research-paper pretraining mix. ²⁵

Latency claim: plausible only with caveats. The TTM block should be fast. The unproven part is the **local LLM narration** and document-processing path on commodity hardware, because IBM’s official Granite 4.1 materials do not publish M2/4060 local throughput numbers. If the memo promises “under 60 seconds,” it should be framed as a **target**, not as a validated performance number. ²⁶

The most defensible rewrite of the pitch would be: **“PIT WALL is a race-engineering workflow copilot that combines document extraction, domain-specific telemetry feature engineering, lightweight multivariate forecasting, and LLM explanation. IBM Granite provides attractive building blocks, but**

the competitive moat comes from motorsport data, labels, and workflow integration.” That is materially weaker than the original startup-style claim, but it is much harder to shoot down. ²⁷

Open questions and limitations

Two important uncertainties remain. First, I did not have access to your actual telemetry schema, label definition, or holdout design, so I cannot judge whether “next-session lap pace” is a well-posed supervised target or whether you have enough data to fine-tune robustly. Second, IBM’s official Granite 4.1 materials reviewed here do not publish local latency benchmarks for the exact hardware examples you care about, so the end-to-end 60-second claim remains partly inferential rather than benchmarked. ²⁸

There is also a productization caveat. IBM’s newer **TTM-R3** research card looks substantially more production-oriented on throughput and adds a **multi-quantile forecasting head**, but it is released under a **non-commercial CC-BY-NC-SA** license and the associated paper was still marked “to be released” in the model card I reviewed. That makes it interesting as research evidence, but not a clean commercial anchor for a startup pitch today. ²⁹

¹ ³ ⁴ ¹² ¹³ ¹⁵ ²³ ²⁷ [ibm-granite/granite-timeseries-ttm-r2](https://huggingface.co/ibm-granite/granite-timeseries-ttm-r2) · Hugging Face
<https://huggingface.co/ibm-granite/granite-timeseries-ttm-r2>

² ⁵ ⁶ ⁷ ⁸ ⁹ ¹⁰ ¹¹ ¹⁴ ¹⁶ ²⁵ ²⁶ https://proceedings.neurips.cc/paper_files/paper/2024/file/874a4d89f2d04b4bcf9a2c19545cf040-Paper-Conference.pdf
https://proceedings.neurips.cc/paper_files/paper/2024/file/874a4d89f2d04b4bcf9a2c19545cf040-Paper-Conference.pdf

¹⁷ ²⁸ <https://huggingface.co/ibm-granite/granite-4.1-8b>
<https://huggingface.co/ibm-granite/granite-4.1-8b>

¹⁸ <https://ollama.com/library/granite4.1>
<https://ollama.com/library/granite4.1>

¹⁹ <https://www.ibm.com/granite/docs/models/granite4-1>
<https://www.ibm.com/granite/docs/models/granite4-1>

²⁰ ²² ²⁴ <https://www.sciencedirect.com/science/article/pii/S2451958824000472>
<https://www.sciencedirect.com/science/article/pii/S2451958824000472>

²¹ <https://arxiv.org/abs/2501.07034>
<https://arxiv.org/abs/2501.07034>

²⁹ <https://huggingface.co/ibm-research/ttm-r3>
<https://huggingface.co/ibm-research/ttm-r3>