

APEX Feature Judging-Rubric Lift: Rank-Ordered Analysis

IBM SkillsBuild AI Builders Challenge — May 2026 Submission

Executive Summary

The IBM SkillsBuild AI Builders Challenge evaluates submissions across four axes: **Technical Execution**, **Innovation**, **Challenge Fit**, and **Implementation & Feasibility**. Within those axes, BeMyApp competitions historically weight live, interactive demos heavily — judges score projects on a 1–5 scale across these criteria and the submission includes a 3-minute demo video plus a public GitHub repository. Published hackathon judging literature, UX research on interactive demos, and BeMyApp's own prize structure converge to reveal a predictable hierarchy of feature impact. The 13 APEX features are ranked below from highest expected rubric lift to lowest, with evidence for each.^{[1][2]}

The Judging Rubric: Axis Weights in Context

The official IBM SkillsBuild / BeMyApp judging rubric for the May 2026 challenge lists four criteria:

Axis	Description
Technical Execution	Effective use of IBM and open source technologies, functional + well-structured solution
Innovation	Creativity, originality, unique application of AI
Challenge Fit	Relevance to racing/F1 context, addressing real-world problems
Implementation & Feasibility	Practicality, scalability, real-world potential

^{^2}

In addition, the challenge has four prize sub-categories: 1st Place (\$2,250), Runner-Up (\$1,250), **Most Innovative** (\$750), and **Best Use of Technology** (\$750). A Grand Prize (\$5,000) is awarded "based on technical execution, innovation, challenge fit, and real-world impact."^[^2]

This prize architecture means features should be evaluated not just against the primary rubric, but against whether they move the needle on specific prize tracks — a key insight from Devpost's analysis of sponsor prize alignment.^[^3]

IBM itself ranks **transparency** as a paramount value for its AI ecosystem. IBM Granite earned the highest-ever score (95%) on Stanford's Foundation Model Transparency Index in 2025,

and IBM's challenge narrative explicitly calls for AI solutions that "drive smarter decisions, improve understanding, and build trust through explainability." This directly inflates rubric scores for any feature demonstrating explainability, transparency, and verifiability.^{[4][5][^2]}

Feature Rank-Order by Expected Judging-Rubric Lift

Rank 1 — (e) Per-Track Sponsor Tightening Copy: 4 BeMyApp Prize Categories with "Wins X Because Y" Explicit Claims

Expected lift: Highest / Prize-Multiplier

Explicit claim alignment to named prize categories is the single highest-leverage move in any BeMyApp hackathon. Devpost's official best-practices guide for hackathon organizers describes the prize structure as a "contract between you and the public" — judges evaluating sponsor tracks need to see projects explicitly demonstrate eligibility, or risk being passed over even if technically superior. The MLH Hackathon Organizer Guide confirms that "sponsored prize tracks significantly boost interest" but require the team to "bring a representative" who understands the challenge, and submissions should make prize-track alignment explicit.^{[6][3]}

For APEX specifically, the BeMyApp challenge has four prize categories: 1st Place, Runner-Up, Most Innovative, and Best Use of Technology. Having per-track copy that says "Wins *Most Innovative* because [specific innovation]" and "Wins *Best Use of Technology* because [IBM Granite/LangFlow usage]" directly reduces judge cognitive load and removes guessing from the scoring process. Published hackathon post-mortems (e.g., Bizthon's analysis of 10 winning traits in 2025 hackathons) show that "judges can smell a forced use case" but equally reward projects that make their prize eligibility legible. This is purely a presentation/packaging feature with zero implementation risk — it adds value regardless of what else the submission contains.^{[7][2]}

Rank 2 — (d) Real Granite TTM In-Browser via Transformers.js (Lazy-Load + WebGPU; Opt-In)

Expected lift: Very High / Technical Execution + Innovation + Best Use of Technology

This feature fires on all four judging axes simultaneously. Technically: running IBM's own Granite TTM (TinyTimeMixer) model live in the browser via Transformers.js v4 + WebGPU is state-of-the-art. Hugging Face released Transformers.js v4 with WebGPU support in March 2026, and the Hacker News community called the combination "a game changer" for browser-based AI. IBM's own developer tutorials highlight TTM as the flagship time-series foundation model for forecasting applications.^{[8][9][10][11][^12]}

For the "Best Use of Technology" prize, the use of a named IBM technology (Granite) in a novel deployment mode (client-side WebGPU, zero server cost) directly satisfies the criterion. The opt-in + lazy-load architecture addresses the privacy argument that WebGPU proponents cite — "no telemetry tracking user queries as long as everything remains within

the browser" — and directly maps to the challenge's "build trust through explainability" mandate. The main risk is device compatibility (~70% WebGPU support as of late 2024), which the opt-in model mitigates.^{[9][13][^2]}

IBM's Granite Time Series workshop materials, including NeurIPS 2023 demo materials, are publicly available at GitHub, confirming this is a recognized showcaseable capability in IBM's own AI portfolio. Running TTM in-browser is a step beyond what most hackathon competitors will attempt — it signals genuine technical depth.^[^14]

Rank 3 — (l) LangGraph Runtime Orchestration Trace at /judges (6-Node State-Machine Visualization)

Expected lift: Very High / Technical Execution + Innovation + Explainability

LangGraph is explicitly listed as one of the required IBM-endorsed technologies in the challenge ("Langflow" is listed; LangGraph from LangChain is the de-facto standard for state-machine agent orchestration). Arize AI's 2024 deep-dive on LangGraph identifies its visual state-machine trace as the key feature for "monitoring agent performance and identifying bottlenecks," and notes that "visual debugging" is a core enterprise-grade feature of the framework. The 2025 LangGraph State-Machine review calls tracing and visualization "a requirement for production AI."^{[15][16][^2]}

IEEE VIS's 2025 session on "Visualization for Explainable AI" and the VISxAI workshop explicitly frame interactive AI system visualizations as the primary mechanism for building judge/stakeholder trust in AI decision-making pipelines. A 6-node state-machine visualization on /judges makes the orchestration architecture self-documenting. For IBM-affiliated judges who understand watsonx Orchestrate and agentic AI, this is the technical depth signal they are trained to recognize — the Agentic AI Hackathon with IBM watsonx Orchestrate explicitly required submissions to "clearly demonstrate how IBM watsonx Orchestrate is used."^{[17][18][^19]}

Rank 4 — (b) Real-Time Apex-Cam Pipeline Visualization (6-Stage Cycling at 1.5s/Stage)

Expected lift: High / Technical Execution + Challenge Fit + Visual Storytelling

Animated, real-time pipeline visualizations have strong UX research support. A 2025 controlled experiment (n=240) in *Nature Scientific Reports* found that live Conversational Agents ($\beta=5.32$, $p<0.001$) and Adaptive Feedback Mechanisms ($\beta=4.86$, $p<0.001$) had the highest positive impact on system usability among interactive interface features in AI-driven systems. The principle extends directly to animated pipeline stages: they reduce cognitive load by making invisible processes visible.^[^20]

Dev3lop's published analysis of animated data visualization best practices identifies pipeline flow animations as particularly effective at "converting technical expertise into accessible insights for broad stakeholder comprehension." The 200–500ms transition range is cited as optimal for comprehension without distraction — the 1.5s stage duration for APEX falls

within a healthy range. The SAFEXPLAIN project (BSC/EU), which targets certifiable explainable AI for critical systems, uses a Core Demo with stage-by-stage animated pipeline visualization as its central communication device for regulators and judges.^{[21][22][23]}

For F1-specific challenge fit, IBM's own partnership with Ferrari uses "real-time telemetry insights" as the flagship use case for AI in racing. A 6-stage cycling visualization of an apex detection pipeline aligns directly with the May Challenge's "AI Perception & Vision" and "AI Analytics" suggested categories.^{[24][25][2]}

Rank 5 — (g) Multi-Driver /compare View (Baseline-vs-Improved Side-by-Side with Delta Visualization)

Expected lift: High / Challenge Fit + Technical Execution

Side-by-side comparison visualizations have strong empirical support. UXmatters' research on comparative decision-making establishes that "we determine value by comparing and contrasting one thing to another — people do not make judgments in a vacuum." MeasuringU's UX research on difference-score visualizations finds that "graphs of difference scores are a helpful visualization technique to highlight differences compared to displaying raw scores." ACM DL's 2024 study on interaction techniques for comparison confirms that "juxtaposition arranges different visual entities in a side-by-side layout, in order to allow users to see similarities and differences across the visualizations."^{[26][27][28]}

In the F1 telemetry domain, head-to-head driver comparison is the canonical power-user feature. F1 Live Pulse and F1 Matcha both offer "head-to-head lap comparison" with side-by-side breakdowns as their headline analytic capability. An F1 driver performance analyzer on GitHub (May 2025) explicitly implements "Driver vs. Driver: a head-to-head lap time comparison" as the primary showcase feature. The delta visualization makes improvement quantifiable in a single glance — directly addressing the challenge's mandate to "turn complex telemetry into clear insights." Hackathon UX research consistently shows that judges spend 60–120 seconds asking "Does it actually work? Is this impressive?" — delta visualizations answer this efficiently.^{[29][30][31][32][2]}

Rank 6 — (a) Interactive COA Gate Toggle (/judges: Live Brake-Throttle Simultaneity Verdict Flip)

Expected lift: High / Innovation + Explainability + Challenge Fit

Interactive counterfactual / toggle interfaces have strong backing in the XAI literature. A 2025 SCITEPRESS study on "Designing Explainable and Counterfactual-Based AI Interfaces" found that operators preferred combinations of counterfactual and feature-importance explanations for understanding AI decisions. The SAFEXPLAIN project's core demo highlights "safety pattern scenarios — where AI influences, but does not solely determine, decision-making" as the ideal XAI demonstration format for critical systems.^{[33][21]}

The COA gate toggle maps precisely onto the IBM challenge's stated goal: "AI can drive smarter decisions, improve understanding, and build trust through explainability." A live

toggle that flips a verdict from "permitted" to "blocked" while showing the brake-throttle simultaneity logic is a textbook interactive XAI demo. IEEE VIS's 9th Workshop on Visualization for AI Explainability calls for "explainables / explorables" — interactive documents that explain how AI techniques work — as the gold standard for AI communication. This feature is the closest APEX has to a purpose-built explainable AI interface, and IBM's judging culture, shaped by the company's #1 transparency ranking from Stanford, would reward it.^{[5][4][17][2]}

Rank 7 — (j) NeurIPS Workshop Paper Draft Included in Submission Package

Expected lift: Moderate-High / Innovation + Technical Depth Signal

Including a NeurIPS workshop paper draft signals academic rigor that most hackathon submissions cannot match. A Reddit thread on the ML community confirms that NeurIPS workshop submissions are becoming "close to being as competitive as some top conferences," with acceptance rates dropping to ~20% as of 2025. NeurIPS workshops are not archival, meaning an early-stage draft is an appropriate and legitimate submission artifact. The NeurIPS 2022 Trustworthy and Socially Responsible ML Workshop explicitly includes "explainable and interpretable AI" and "verifiable guarantees" as submission criteria — directly aligned with APEX's core proposition.^{[34][35][36]}

However, the lift here is bounded: the IBM SkillsBuild rubric does not explicitly weight academic publications, and judges reviewing hundreds of submissions may not allocate meaningful attention to a PDF draft. The Mozilla Foundation's analysis of EleutherAI's Evaluation Harness describes how academic framing adds credibility ("one of this year's Mozilla Technology Fund awardees") but converts to prize recognition only when paired with a working demo. The paper functions best as a trust amplifier for the /methodology and /lips-harness pages rather than as a standalone scoring artifact.^{[37][2]}

Rank 8 — (k) APEX-Bench Evaluation Harness at /lips-harness (4-Axis Ablation: Latency + Integrity + Physics + Skill)

Expected lift: Moderate-High / Technical Execution + Innovation

A structured evaluation harness is a strong technical credibility signal. EleutherAI's LM Evaluation Harness, highlighted by Mozilla's Technology Fund, demonstrates that "prior to the Evaluation Harness, developers were creating models, evaluating them according to their own code, and not releasing the results — making testing difficult and tedious." Publishing APEX-Bench at /lips-harness directly addresses reproducibility — a value that IBM's transparency-first culture prizes.^{[4][37]}

A 2026 CMU Library hackathon on reproducibility explicitly framed the challenge as "can researchers replicate and augment published research, working with existing data and code?" The DoraHacks judging guide for hackathons describes AI evaluation as adding value by "applying criteria consistently without fatigue" and providing "transparent reasoning."

However, the 4-axis ablation study format is best appreciated by technical judges with ML evaluation experience — its impact may be uneven across a multi-judge panel.^[38]^[39]

Rank 9 — (f) /judge-tour Walkthrough Mode (6-Step Narrative Walkthrough)

Expected lift: Moderate / Polish + Documentation + Usability

Product tour walkthroughs reduce judge friction, which is consistently cited in hackathon post-mortems as a non-trivial predictor of evaluation quality. The Opportunity Hack 2025 judging guide identifies "documentation" as one of four primary judging criteria, asking: "Can people use the product as it stands today? Does it have enough documentation? Is it easy to use?" DoraHacks' analysis advises that "every bit of friction you remove is attention judges can redirect to the projects themselves."^[40]^[38]

Appcues' 2026 guide on product tours confirms they "accelerate time-to-value, reduce user confusion, and improve feature discovery." Userflow's research identifies product tours as effective for "boosting activation" and "reducing support requests." A 6-step narrative walkthrough for /judges is a well-understood UX pattern with evidence-backed lift in comprehension and satisfaction. The constraint is that the IBM challenge judges review submissions asynchronously via a 3-minute video + GitHub — a walkthrough on a live page adds secondary value only if judges visit the deployed site.^[41]^[42] ^[^2]

Rank 10 — (i) Engine-Agnostic Byte-Equality Demo at /judges (Side-by-Side V1+V2 Violation Strings Byte-Identical Modulo ENGINE)

Expected lift: Moderate / Technical Execution + Trustworthiness

Byte-level reproducibility and determinism are emerging as differentiators in AI evaluation, but their appreciation requires technical sophistication from judges. A 2026 arXiv paper on EigenAI states that "deterministic inference forms the technical cornerstone of EigenAI's verifiability framework. Without strict bit-level reproducibility, there is no basis for verifiable results." The LinkedIn analysis from a senior enterprise AI consultant (2025) frames the debate clearly: "confidence in high-stakes domains doesn't come from seeing inside the model — it comes from knowing it will produce the right outcome, every time. That's determinism."^[43]^[44]

For the F1 motorsport context, where rule enforcement (COA violations) requires consistent, auditable verdicts across different race weekends, byte-equality is a meaningful engineering claim. However, hackathon rubrics rarely include a "determinism" axis explicitly. The Bellingcat hackathon rubric, for instance, explicitly excludes "code quality" and "architecture/use of software design patterns" from judging. This feature scores well on technical depth but is unlikely to independently move rubric needle unless a judge with motorsport or safety-critical AI background evaluates it.^[45]^[46] ^[^47]

Rank 11 — (c) /changelog Auto-Rendered from Git Log (Server Component Reading Git Log; Conventional-Commit Color Pills)

Expected lift: Moderate-Low / Polish + Documentation

Conventional Commits and auto-generated changelogs are respected developer practices. The conventional commits specification is designed to produce "human and machine-readable" commit histories and enable automated changelog generation. Hackathon guides consistently identify documentation as a judging criterion — Opportunity Hack scores it explicitly.^[48]^[40]

However, published hackathon rubrics draw a clear distinction between internal developer practice (commit quality) and user-facing value. The Hacker News discussion of conventional commits notes that "the real problem is that commit messages and a changelog serve 2 different purposes and have different audiences." For an IBM hackathon with a 3-minute demo video as the primary evaluation artifact, a /changelog page — however elegantly rendered — is a secondary signal. It contributes polish and signals engineering discipline, but does not independently score on any of the four judging axes.^[49]^[2]

Rank 12 — (h) Sookra Methodology Five Pillars Landing at /methodology

Expected lift: Low-Moderate / Challenge Fit (contextual)

A dedicated methodology page demonstrates intellectual rigor and structured thinking, which maps loosely onto the "challenge fit" and "feasibility" axes. CMU Energy Hackathon rubrics ask whether "recommendations are backed by sound economic reasoning." Having a named methodology (Sookra Five Pillars) implies a principled approach.^[50]

However, this feature's impact is entirely dependent on the methodology's actual content and how legibly it maps to IBM's challenge mandate. Without the specific pillars being directly referenced in the 3-minute demo video or README — both of which are the primary review artifacts — a standalone /methodology page functions primarily as a trust amplifier for judges who visit the deployed site. It adds credibility but generates minimal independent scoring lift compared to interactive features.^[2]

Rank 13 — (m) OG Cards on Every Route (Judges Sharing URLs See Editorial-Paddock Identity)

Expected lift: Low / Secondary Effect

Open Graph cards are universal best practice for any deployed web application. The OG protocol, created by Facebook in 2010, is now supported by LinkedIn, Slack, Discord, WhatsApp, and virtually every platform that renders link previews. Links with compelling OG preview cards get "2–3× more clicks than bare URL strings." For social media engagement and brand identity, OG cards are table stakes.^[51]^[52]^[53]

However, in the context of hackathon judging, OG cards operate at the margin. Judges receive submissions through the BeMyApp platform and evaluate them against the rubric — they are not browsing a Twitter/X feed where rich link previews drive discovery. The value accrues if judges share specific routes on Slack or Discord during deliberation, creating an indirect identity signal. This is real but not measurable against the rubric axes. Several hackathon guides explicitly list "non-functional visuals and aesthetics" and "pretty/polished frontend" as non-criteria (Bellingcat explicitly excludes them). OG cards are polish — they should be implemented because they signal professionalism, but they do not independently elevate judging scores.^{[47][2]}

Consolidated Rank Table

Rank	Feature	Primary Rubric Axes	Evidence Basis	Lift Level
1	(e) Per-track sponsor prize copy	All 4 + Prize Multiplier	BeMyApp/Devpost prize alignment research ^{[3][6]}	★★★★★
2	(d) Granite TTM in-browser / WebGPU	Technical Execution + Innovation + Best Use of Tech	IBM Granite TTM ^{[12][14]} , Transformers.js v4 ^[^9] , WebGPU ^[^13]	★★★★★
3	(l) LangGraph orchestration trace	Technical Execution + Innovation + Explainability	LangChain/Arize ^{[15][16]} , IBM watsonx Orchestrate ^[^18] , VISxAI ^[^17]	★★★★½
4	(b) Real-time pipeline visualization	Technical Execution + Challenge Fit	Nature IIF study ^[^20] , Dev3lop ^[^23] , SAFEXPLAIN ^[^21]	★★★★
5	(g) Multi-driver /compare + delta	Challenge Fit + Technical Execution	UXmatters ^[^26] , MeasuringU ^[^27] , F1 Live Pulse ^{[30][31]}	★★★★
6	(a) COA gate toggle interactive	Innovation + Explainability + Challenge Fit	XAI interfaces ^[^33] , SAFEXPLAIN ^[^21] , VISxAI ^[^17]	★★★★½
7	(j) NeurIPS workshop paper draft	Innovation + Technical Depth Signal	NeurIPS workshop competitiveness ^{[34][35]} , Mozilla Eval Harness ^[^37]	★★★
8	(k) APEX-Bench evaluation harness	Technical Execution + Innovation	EleutherAI/Mozilla Harness ^[^37] , CMU Reproducibility ^[^39] , DoraHacks ^[^38]	★★★
9	(f) /judge-tour walkthrough	Polish + Documentation + Usability	Opportunity Hack ^[^40] , Appcues ^[^42] , DoraHacks ^[^38]	★★½
10	(i) Byte-equality engine-agnostic demo	Technical Execution (niche)	EigenAI arXiv ^[^44] , determinism research ^[^43]	★★
11	(c) /changelog from git log	Polish + Documentation	Conventional commits ^[^48] , Hackathon doc criteria ^[^40]	★★
12	(h) Sookra Methodology Five Pillars	Challenge Fit (contextual)	CMU rubric ^[^50] , hackathon feasibility criteria	★½

Rank	Feature	Primary Rubric Axes	Evidence Basis	Lift Level
13	(m) OG cards on every route	Secondary / Identity Signal	OG protocol ^[51] [52][[^] 53], Bellingcat rubric ^{[[^]47]}	★

Cross-Cutting Analysis

The IBM Transparency Premium

IBM's corporate positioning around Granite's #1 transparency score (95% FMTI, Stanford 2025) and the challenge's explicit mandate for "explainability" and "trust" means that any feature demonstrating visible, interactive AI decision-making gets a multiplier on the Innovation and Challenge Fit axes. Features (d), (l), (a), and (b) all directly embody this — they make the AI's behavior visible, inspectable, and trustworthy in real-time. IBM and Ferrari's 2026 partnership explicitly showcases "real-time AI insights" as the flagship use case for AI in motorsport.^[25][5][^[24][4][[^]2]

The Live Demo Premium

A 2026 YouTube session by a judge who reviewed 21 hackathon pitches in one day found that "half couldn't even show their demo" — making a functioning, accessible deployed demo one of the strongest differentiators. The IBM challenge requires a 3-minute demo video, but a deployed site with interactive features extends the evaluation window beyond what video allows. Features accessible live at named routes (/judges, /compare, /methodology) maximize this advantage.^[54][2]

The Best Use of Technology Pathway

Features (d) and (l) are the clearest pathways to the "Best Use of Technology" prize (\$750). IBM Granite TTM is a named IBM technology used in a novel deployment mode; LangGraph is the de-facto standard for the agentic AI workflows IBM's challenge promotes. The prize requires demonstrating IBM technology use — these two features make that explicit and demonstrable.^{[[^]2]}

Diminishing Returns in the Lower Tier

Features (c), (h), and (m) are professional engineering practices that contribute to polish and documentation scores, which the Opportunity Hack rubric weights but most BeMyApp rubrics do not list as standalone criteria. They are worth shipping for completeness but should not consume sprint cycles at the expense of the top-tier features.^[40][2]

Caveats and Limitations

The IBM SkillsBuild AI Builders Challenge does not publish quantitative axis weights — the rubric lists four criteria without percentage allocations. This analysis infers weighting from: (1) the prize category structure (Most Innovative vs. Best Use of Technology as separate prizes suggests equal weighting on both Innovation and Technical Execution); (2) IBM's corporate AI positioning around transparency and Granite; (3) comparable BeMyApp/hackathon rubric structures where available. No published post-mortem exists specifically for this challenge series, as it launched in May 2026. Rankings for features in the upper tier are supported by multiple independent evidence sources; rankings in the lower tier reflect best-judgment synthesis.^[55][2]

References

1. [IBM SkillsBuild AI Hackathon at OSU - BeMyApp](#) - AI Hackathon at Ohio State is a 24-hour, in person student hackathon powered by IBM SkillsBuild. Stu...
2. [AI Builders Challenge, sponsored by IBM SkillsBuild, for Students | Build AI & Win Prizes](#) - Calling all university students. Build real-world AI projects using IBM technologies in two innovati...
3. [Best Practices: Incorporating Hackathon Sponsors](#) - Featuring sponsors is a great way to include partners, add to the prize pool, and increase the marke...
4. [IBM recognized as the leader in AI transparency](#) - IBM's leadership in the Foundation Model Transparency Index reflects our philosophy that organizatio...
5. [IBM Granite is ranked world's most transparent model](#) - The Stanford University Foundation Model Transparency Index has ranked IBM Granite number one this y...
6. [Potential Sponsor Perks - Hackathon Organizer Guide](#)
7. [10 Winning Hacks: What Makes a Hackathon Project Stand ...](#) - You've seen them at every hackathon — projects that make judges lean forward, teams that dominate th...
8. [WebGPU and Transformers.js = Local LLMs in the Browser](#) - I wrote up a deep dive into how it works, what WebGPU enables, and the future of local browser-based...
9. [Huggingface has just released Transformer.js v4 with WebGPU ...](#) - Transformers.js allows you to run models right in the browser. The fourth version focuses on perform...
10. [Why WebGPU + Transformers.js is a Game Changer - YouTube](#) - Transformers.js is the perfect way to run AI models directly in browsers. This library from HuggingF...
11. [Zero-Shot Time Series Forecasting Using Granite TTM on OCI](#) - To get started, open AI Quick Actions from the Data Science Notebook and navigate to the Model Explo...
12. [Using the IBM Granite models for time series forecasting](#) - In this tutorial, you'll learn how to use the TinyTimeMixer (TTM) model, a compact time-series found...
13. [Running models on WebGPU - Hugging Face](#) - WebGPU is a new web standard for accelerated graphics and compute. The API enables web developers to...

14. [GitHub - ibm-granite/granite-tsfm: Foundation Models for Time Series](#) - Run the TTM tutorial in Google Colab, and quickly build a forecasting application with the pre-train...
15. [Tracing and Evaluating LangGraph Agents - Arize AI](#) - In this post, we'll discuss how LangGraph's traces can be ingested into Arize, and how to leverage L...
16. [LangGraph 2025 Review: State-Machine Agents for Production AI](#) - You can also visualize state machines as graphs, where: States are represented as nodes. Transitions...
17. [9th Workshop on Visualization for AI Explainability](#) - 9th Workshop on. Visualization for AI Explainability. October 2026 at IEEE VIS in Boston, MA. VISxAI...
18. [Agentic AI Hackathon with IBM watsonx Orchestrate - Lablab.ai](#) - Explore the recap of Agentic AI Hackathon with IBM watsonx Orchestrate. Browse winning projects, sub...
19. [Visualization for Explainable AI \(XAI\) - IEEE VIS](#) - Monday, Nov 3rd, 2025 @ 07:00 – 08:30. Chair: Roy A. Ruddle, Liqun Liu, Mohsen Azarmi. 0 presentatio...
20. [Modelling the impact of interactive interface features on user ...](#) - This paper addresses these knowledge gaps by developing a comprehensive model that quantifies the im...
21. [BSC's SAFEXPLAIN Project: 'Toward Explainable AI in Critical ...](#) - Barcelona, 03 July 2025 -- The SAFEXPLAIN project, formed by an inter-disciplinary consortium of six...
22. [Making certifiable AI a reality for critical systems: SAFEXPLAIN core demo](#) - AI is taking on a dominant role in critical systems, powering new autonomous functions. However, cur...
23. [Animated Transitions in Interactive Data Visualizations - dev3lop](#) - Animated transitions bring significant value: they guide users effortlessly between distinct visuals...
24. [Ferrari is using IBM's AI to create F1 superfans - TechCrunch](#) - The IBM-Ferrari partnership centers on that idea of storytelling, enhancing fan engagement by overha...
25. [IBM powers real-time insights for Ferrari F1 fans with AI and data ...](#) - The Ferrari mobile app has been reimagined using AI and analytics powered by watsonx, turning comple...
26. [The Power of Comparison: How It Affects Decision Making - UXmatters](#) - We determine value by comparing and contrasting one thing to another. People do not make judgments a...
27. [Visualizing Data: Raw vs Difference Scores - MeasuringU](#) - Graphs of difference scores are a helpful visualization technique to help highlight differences comp...
28. [Interaction Techniques for Comparing Video - ACM Digital Library](#) - Juxtaposition arranges different visual entities in a side-by-side layout, in order to allow users t...
29. [Pitch Perfect: How to Present Your Hack Like a Pro - Medium](#) - Your smart contracts are deployed. Your frontend works flawlessly. Your documentation is pristine. Y...
30. [Formula Live Pulse | Your Real-time F1 Companion](#) - Get real-time Formula 1 data, telemetry, and insights with Formula Live Pulse. Track live timing, an...
31. [GitHub - rembertdesigns/f1-driver-performance-analyzer: An interactive dashboard using Streamlit and Plotly to analyze and compare F1 driver performance, featuring head-to-head stats, stint analysis, and an AI scoring model.](#) - An interactive dashboard using Streamlit and Plotly to analyze and compare F1 driver performance, fe...

32. [Visualize OpenF1 Data | F1 2025 - F1 Telemetry](#) - Explore F1 driver performance for 2025 with live telemetry and real race data. Compare drivers, anal...
33. [Designing Explainable and Counterfactual-Based AI Interfaces for ...](#)
34. [Are Neurips workshop competitive? \[R\] : r/MachineLearning - Reddit](#) - This year it seems that we are getting ~500 submissions, compared to last year's ~200. If ~100 are a...
35. [Machine Learning and the Physical Sciences, NeurIPS 2025](#) - Neither the paper checklist nor a broader impact statement are required for workshop submissions. If...
36. [Workshop at NeurIPS 2022](#) - Trustworthy and Socially Responsible Machine Learning | Workshop at NeurIPS 2022
37. [Evaluation Harness Is Setting the Benchmark for Auditing Large ...](#) - Evaluation Harness is an open-source tool that allows researchers and developers to put their LLMs t...
38. [Use Ai To Handle Volume, Or...](#) - A Practical Guide for Organizers Your hackathon went very well. Submissions are in. Now comes the...
39. [2025 Reproducibility Hackathon Investigates LLM Writing Styles](#) - 35 participants were challenged to replicate and augment published research, working with existing d...
40. [Hackathon Judge: Opportunity Hack Intro 2025 - YouTube](#) - Learn more about judging criteria for Opportunity Hack 2025 as we hack at ASU in Tempe, Arizona Octo...
41. [How To Build Product Tours That Boost SaaS Success - Userflow](#) - This guide will cover everything you need to know about product tours, from their benefits and types...
42. [Product tour guide: boost activation and retention - Appcues](#) - Build a product walkthrough that drives activation: a 6-step framework, UI patterns, best practices,...
43. [Determinism vs Transparency: Which AI do we want? - LinkedIn](#) - Where does confidence in AI really come from—transparency, or determinism? Evaluation frameworks lea...
44. [EigenAI: Deterministic Inference, Verifiable Results - arXiv](#) - Deterministic inference forms the technical cornerstone of EigenAI's verifiability framework. Withou...
45. [JUDGING RUBRIC - Ethical AI Hackathon - CGU Research Centers](#) - Instructions for Judges: Review each category below and use the score guide (1 = Poor, 5 = Excellent...
46. [\[PDF\] 2025 CMBB Biotech Bridge Hackathon Beat the Clock: Innovating ...](#)
47. [\[PDF\] Bellingcat Hackathons Grading Rubric](#)
48. [A specification for adding human and ... - Conventional Commits](#) - A specification for adding human and machine readable meaning to commit messages
49. [I work with projects where this is utilized and I'm not sure they are ...](#)
50. [CMU Energy + Cleantech Hackathon - Judging Criteria](#) - Judging criteria for the CMU Energy + Cleantech Hackathon
51. [Common Mistakes To Avoid](#) - Master Open Graph tags to control how your content appears on Facebook, Twitter, and LinkedIn. Inclu...
52. [Open Graph and Meta Tags: Controlling How Your Content Appears ...](#) - Control how your pages look when shared on social media. Master Open Graph tags, Twitter Cards, imag...

53. Open Graph Protocol Explained: How OG Tags Work (2026) - Preview and debug Open Graph tags, Twitter Cards, and social media previews. Fix broken previews bef...
54. I Judged 21 Hackathon Pitches. Half Couldn't Even Show Their Demo - How do you pitch your idea in under 10 minutes and actually leave time for the demo? I recently judg...
55. IBM SkillsBuild AI Builders Challenge - Eventbrite - Monthly Student Challenges & Learning Labs | \$15,000 Prize Pool